

# Introduction to SQL – 2

Debapriyo Majumdar

DBMS – Fall 2016

Indian Statistical Institute Kolkata

*Slides re-used, with minor modification, from*

*Silberschatz, Korth and Sudarshan*

[www.db-book.com](http://www.db-book.com)

# Outline

---

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions
- Nested Subqueries
- Modification of the Database

# Nested Subqueries

---

- A **subquery** is a **select-from-where** expression that is nested within another query
- The nesting can be done in the SQL query

**select**  $A_1, A_2, \dots, A_n$   
**from**  $r_1, r_2, \dots, r_m$   
**where**  $P$

as follows:

- $A_i$  can be replaced by a subquery that generates a single value
- $r_i$  can be replaced by any valid subquery
- $P$  can be replaced with an expression of the form:

$B <\text{operation}> (\text{subquery})$

Where  $B$  is an attribute and  $<\text{operation}>$  to be defined later

# Subqueries in the Where Clause

---

- A common use of subqueries is to perform tests
  - For set membership
  - For set comparisons
  - For set cardinality

# Set Membership

---

- Find courses offered in Fall 2009 **and** in Spring 2010

```
select distinct course_id  
from section  
where semester = 'Fall' and year = 2009 and  
       course_id in (select course_id  
                       from section  
                       where semester = 'Spring' and year = 2010);
```

- Find courses offered in Fall 2009 **but not in** Spring 2010

```
select distinct course_id  
from section  
where semester = 'Fall' and year = 2009 and  
       course_id not in (select course_id  
                           from section  
                           where semester = 'Spring' and year = 2010);
```

# Set Membership (Continued)

---

- Find the total number of (distinct) students who have taken course sections taught by the instructor with *ID* 10101

```
select count (distinct ID)  
from takes  
where (course_id, sec_id, semester, year) in  
      (select course_id, sec_id, semester, year  
       from teaches  
       where teaches.ID = 10101);
```

**Note:** The above query can be written in a much simpler manner.  
The formulation above is simply to illustrate SQL features.

# Set Comparison – “some” Clause

---

- Find names of instructors with salary greater than that of some (at least one) instructor in the Biology department

```
select distinct T.name  
from instructor as T, instructor as S  
where T.salary > S.salary and S.dept name = 'Biology';
```

- Same query using > **some** clause

```
select name  
from instructor  
where salary > some (select salary  
                        from instructor  
                        where dept name = 'Biology');
```

# Definition of “some” Clause

---

- $F \text{ <comp> some } r \Leftrightarrow \exists t \in r \text{ such that } (F \text{ <comp> } t)$   
Where <comp> can be: <, ≤, >, =, ≠

$$(5 \text{ < some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{true} \quad (\text{read: } 5 \text{ < some tuple in the relation})$$

$$(5 \text{ < some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{false}$$

$$(5 = \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true}$$

$$(5 \neq \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true} \quad (\text{since } 0 \neq 5)$$

(= some) ≡ in

However, (≠ some) ≠ not in



# Set Comparison – “all” Clause

---

- Find the names of all instructors whose salary is greater than the salary of all instructors in the Biology department.

```
select name  
from instructor  
where salary > all (select salary  
                        from instructor  
                        where dept name = 'Biology');
```

# Definition of “all” Clause

---

- $F \text{ <comp> all } r \Leftrightarrow \forall t \in r (F \text{ <comp> } t)$

$$(5 \text{ < all } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{false}$$

$$(5 \text{ < all } \begin{array}{|c|} \hline 6 \\ \hline 10 \\ \hline \end{array}) = \text{true}$$

$$(5 \text{ = all } \begin{array}{|c|} \hline 4 \\ \hline 5 \\ \hline \end{array}) = \text{false}$$

$$(5 \text{ ≠ all } \begin{array}{|c|} \hline 4 \\ \hline 6 \\ \hline \end{array}) = \text{true (since } 5 \neq 4 \text{ and } 5 \neq 6)$$

$(\neq \text{ all}) \equiv \text{not in}$

However,  $(= \text{ all}) \not\equiv \text{in}$

# Test for Empty Relations

---

- The **exists** construct returns the value **true** if the argument subquery is nonempty.
- **exists**  $r \Leftrightarrow r \neq \emptyset$
- **not exists**  $r \Leftrightarrow r = \emptyset$

# Use of “exists” Clause

---

- Yet another way of specifying the query “Find all courses taught in both the Fall 2009 semester and in the Spring 2010 semester”

```
select course_id
from section as S
where semester = 'Fall' and year = 2009 and
      exists (select *
              from section as T
              where semester = 'Spring' and year = 2010
                  and S.course_id = T.course_id);
```

- **Correlation name** – variable *S* in the outer query
- **Correlated subquery** – the inner query

# Use of “not exists” Clause

---

- Find all students who have taken all courses offered in the Biology department

```
select distinct S.ID, S.name from student as S  
  where not exists ( (select course_id from course  
                    where dept_name = 'Biology')  
                    except  
                    (select T.course_id from takes as T  
                    where S.ID = T.ID));
```

- First nested query lists all courses offered in Biology
  - Second nested query lists all courses a particular student took
- 
- Note that  $X - Y = \emptyset \Leftrightarrow X \subseteq Y$
  - *Note:* Cannot write this query using = **all** and its variants

# Test for Absence of Duplicate Tuples

---

- The **unique** construct tests whether a subquery has any duplicate tuples in its result
- The **unique** construct evaluates to “true” if a given subquery contains no duplicates
- Find all courses that were offered at most once in 2009

```
select T.course_id  
from course as T  
where unique (select R.course_id  
                from section as R  
                where T.course_id= R.course_id  
                    and R.year = 2009);
```

# Subqueries in the From Clause

---

- SQL allows a subquery expression to be used in the **from** clause
- Find the average instructors' salaries of those departments where the average salary is greater than \$42,000

```
select dept_name, avg_salary
from (select dept_name, avg (salary) as avg_salary
      from instructor
      group by dept_name)
where avg_salary > 42000;
```

- Note that we do not need to use the **having** clause
- Another way to write above query

```
select dept_name, avg_salary
from (select dept_name, avg (salary)
      from instructor
      group by dept_name) as dept_avg (dept_name, avg_salary)
where avg_salary > 42000;
```

# With Clause

---

- The **with** clause provides a way of defining a temporary relation whose definition is available only to the query in which the **with** clause occurs.
- Find all departments with the maximum budget

```
with max_budget (value) as  
    (select max(budget)  
     from department)  
select department.name  
from department, max_budget  
where department.budget = max_budget.value;
```



# Complex Queries using With Clause

---

- Find all departments where the total salary is greater than the average of the total salary at all departments

```
with dept_total (dept_name, value) as  
    (select dept_name, sum(salary)  
     from instructor  
     group by dept_name),  
dept_total_avg(value) as  
    (select avg(value)  
     from dept_total)  
select dept_name  
from dept_total, dept_total_avg  
where dept_total.value > dept_total_avg.value;
```

# Scalar Subquery

---

- Scalar subquery is one which is used where a single value is expected
- List all departments along with the number of instructors in each department

```
select dept_name,  
        (select count(*)  
         from instructor  
         where department.dept_name =  
instructor.dept_name)  
        as num_instructors  
from department;
```

- Runtime error if subquery returns more than one result tuple

# Modification of the Database

---

- Deletion of tuples from a given relation
- Insertion of new tuples into a given relation
- Updating of values in some tuples in a given relation

# Deletion

---

- Delete all instructors

**delete from** *instructor*

- Delete all instructors from the Finance department

**delete from** *instructor*

**where** *dept\_name* = 'Finance';

- Delete all tuples in the *instructor* relation for those instructors associated with a department located in the Watson building

**delete from** *instructor*

**where** *dept name* **in**

(**select** *dept name* **from** *department*

**where** *building* = 'Watson');

# Deletion (Continued)

---

- Delete all instructors whose salary is less than the average salary of instructors

**delete from** *instructor*  
**where** *salary* < (**select avg** (*salary*)  
                                 **from** *instructor*);

- Problem: as we delete tuples from deposit, the average salary changes
- Solution used in SQL:
  1. First, compute **avg** (*salary*) and find all tuples to delete
  2. Next, delete all tuples found above (without recomputing **avg** or retesting the tuples)

# Insertion

---

- Add a new tuple to *course*

```
insert into course  
values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```

- or equivalently

```
insert into course (course_id, title, dept_name, credits)  
values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```

- Add a new tuple to *student* with *tot\_creds* set to null

```
insert into student  
values ('3003', 'Green', 'Finance', null);
```

# Insertion (Continued)

---

- Add all instructors to the *student* relation with *tot\_creds* set to 0

**insert into** *student*

**select** *ID, name, dept\_name, 0*

**from** *instructor*

- The **select from where** statement is evaluated fully before any of its results are inserted into the relation.
- Otherwise queries like

**insert into** *table1* **select \*** **from** *table1*

would cause problem

# Updates

---

- Increase salaries of instructors whose salary is over \$100,000 by 3%, and all others by a 5%
  - Write two **update** statements:  

```
update instructor
  set salary = salary * 1.03
  where salary > 100000;
update instructor
  set salary = salary * 1.05
  where salary <= 100000;
```
  - The order is important
  - Can be done better using the **case** statement



# Case Statement for Conditional Updates

---

- Same query as before but with case statement

**update** *instructor*

**set** *salary* = **case**

**when** *salary* <= 100000 **then** *salary* \* 1.05

**else** *salary* \* 1.03

**end**

# Updates with Scalar Subqueries

---

- Recompute and update `tot_creds` value for all students

**update** *student S*

**set** *tot\_cred* = (**select** **sum**(*credits*)

**from** *takes, course*

**where** *takes.course\_id* = *course.course\_id* **and**

*S.ID* = *takes.ID* **and** *takes.grade* <> 'F' **and**

*takes.grade* **is not null**);

- Sets *tot\_creds* to null for students who have not taken any course

- Instead of **sum**(*credits*), use:

**case**

**when** **sum**(*credits*) **is not null** **then** **sum**(*credits*)

**else** 0

**end**