

UNIX Domain Protocols

Sanjaya Kumar Jena

Computer Sc. & Engineering
Faculty of Engineering, ITER
SOA, Deemed To Be University
Bhubaneswar-30

Introduction

- ✍ The Unix domain protocols are not an actual protocol suite, but a way of performing client/server communication on a single host using the same API that is used for clients and servers on different hosts.
- ✍ The Unix domain protocols are an alternative to the interprocess communication (IPC) methods, when the client and server are on the same host.
- ✍ Two types of sockets are provided in the Unix domain: stream sockets (similar to TCP) and datagram sockets (similar to UDP).
- ✍ Even though a **raw socket** is also provided, its semantics have never been documented as per information till now.

Reasons for Unix Domain Protocols

- ✎ On Berkeley-derived implementations, Unix domain sockets are often twice as fast as a TCP socket when both peers are on the same host.
- ✎ Unix domain sockets are used when passing **descriptors** between processes on the same host.
- ✎ Implementations of Unix domain sockets provide the client's credentials (user ID and group IDs) to the server, which can provide additional security checking.

Protocol addresses: Unix Domain Protocols

- ✍ The protocol addresses used to identify clients and servers in the Unix domain are **pathnames within the normal filesystem**.
- ✍ These pathnames are not normal Unix files: We cannot read from or write to these files except from a program that has associated the pathname with a Unix domain socket.
- ✍ IPv4 uses a combination of 32-bit addresses and 16-bit port numbers for its protocol addresses.
- ✍ IPv6 uses a combination of 128-bit addresses and 16-bit port numbers for its protocol addresses.

Unix Domain Socket Address Structure

The Unix domain socket address structure is defined in the header `<sys/un.h>`.

```
#include<sys/un.h>

struct sockaddr_un {

    sa_family_t sun_family; /* AF_LOCAL */

    char sun_path[104];     /* null-terminated pathname */

};
```

- ✍ The pathname stored in the `sun_path` array must be null-terminated.
- ✍ The macro `SUN_LEN` is provided and it takes a pointer to a `sockaddr_un` structure and returns the length of the structure, including the number of non-null bytes in the pathname.
- ✍ The **unspecified** address is indicated by a **null string as the pathname**, that is, a structure with `sun_path[0]` equal to 0. This is the Unix domain equivalent of the IPv4 `INADDR_ANY` constant and the IPv6 `IN6ADDR_ANY_INIT` constant.
- ✍ POSIX renames the Unix domain protocols as “local IPC,” to remove the dependence on the Unix OS.
- ✍ The historical constant `AF_UNIX` becomes `AF_LOCAL`.

Unix Domain Socket Address Structure Filling

UNSPECIFIED CASE

```
#include<sys/un.h>
int listenfd;
struct sockaddr_un servaddr;
listenfd = socket(AF_LOCAL, SOCK_STREAM, 0);
bzero(&servaddr, sizeof(servaddr));
servaddr.sun_family = AF_LOCAL;
servaddr.sun_path[0]=0; /* for unspecified address */
```

Unix Domain Socket Address Structure Fillin

UNSPECIFIED CASE

```
#include<sys/un.h>
int listenfd;
struct sockaddr_un servaddr;
listenfd = socket(AF_LOCAL, SOCK_STREAM, 0);
bzero(&servaddr, sizeof(servaddr));
servaddr.sun_family = AF_LOCAL;
servaddr.sun_path[0]=0; /* for unspecified address */
```

SPECIFIED CASE

```
#include<sys/un.h>
int sockfd;
struct sockaddr_un servaddr;
sockfd = socket(AF_LOCAL, SOCK_STREAM, 0);
unlink(argv[1]);
bzero(&servaddr, sizeof(servaddr));
servaddr.sun_family = AF_LOCAL;
strcpy(servaddr.sun_path, argv[1]); /* for specified */
```


bind, listen, getsockname in Unix Domain Socket Address Structure

```
int main(int argc, char **argv){
int sockfd;    socklen_t len;
struct sockaddr_un addr1, addr2;
if (argc != 2){
    printf("usage: unixbind <pathname>");
    return -1;
}
sockfd = Socket(AF_LOCAL, SOCK_STREAM, 0);
unlink(argv[1]); /* OK if this fails */
bzero(&addr1, sizeof(addr1));
addr1.sun_family = AF_LOCAL;
strncpy(addr1.sun_path, argv[1], sizeof(addr1.sun_path) - 1);
bind(sockfd, (struct sockaddr *) &addr1, SUN_LEN(&addr1));
len = sizeof(addr2);
getsockname(sockfd, (struct sockaddr *) &addr2, &len);
printf("bound name = %s, returned len = %d\n", addr2.sun_path
    , len);
listen(sockfd, 5);
exit(0);
}
```

Unix Domain Stream Server

```
int main(int argc, char **argv){
int listenfd, connfd; socklen_t clilen, len;
struct sockaddr_un cliaddr, servaddr;
len=sizeof(struct sockaddr_un);
listenfd = socket(AF_LOCAL, SOCK_STREAM, 0);
//unlink(argv[1]);
bzero(&servaddr, sizeof(servaddr));
servaddr.sun_family = AF_LOCAL;
//strcpy(servaddr.sun_path, argv[1]); // for specified
servaddr.sun_path[0]=0; //for unspecified address
bind(listenfd, (struct sockaddr *)&servaddr, sizeof(servaddr));
getsockname(listenfd, (struct sockaddr *)&servaddr, &len);
printf("bound name for client= %s\n", servaddr.sun_path);
listen(listenfd, 5);
for ( ; ; ) {
    clilen = sizeof(cliaddr);
    connfd = accept(listenfd, (struct sockaddr *)&cliaddr, &
        clilen);
    write(connfd, "YES\n", 4);
    close(connfd);
}
return 0;}
```

Unix Domain Stream Client

```
int main(int argc, char **argv) {
    int listenfd, connfd, cr, n;
    char buf[12];
    socklen_t clilen;
    struct sockaddr_un cliaddr, servaddr;
    listenfd = socket(AF_LOCAL, SOCK_STREAM, 0);
    bzero(&servaddr, sizeof(servaddr));
    servaddr.sun_family = AF_LOCAL;
    strcpy(servaddr.sun_path, argv[1]);
    cr=connect(listenfd, (struct sockaddr *)&servaddr, sizeof(
        servaddr));
    if (cr==0) {
        printf("connect success\n");
    }
    else{
        printf("Unable to Connect:\n");
        return -1;
    }
    n=read(listenfd, buf, 4);
    write(1, buf, n);
    return 0;
}
```

Unix Domain Datagram Server

```
int main(int argc, char **argv) {
int listenfd, connfd, sr;
char buf[20];
socklen_t clilen, len;
struct sockaddr_un cliaddr, servaddr;
len=sizeof(struct sockaddr_un);
listenfd = socket(AF_LOCAL, SOCK_DGRAM, 0);
unlink(argv[1]);
bzero(&servaddr, sizeof(servaddr));
servaddr.sun_family = AF_LOCAL;
strcpy(servaddr.sun_path, argv[1]); // for specified
bind(listenfd, (struct sockaddr *)&servaddr, sizeof(servaddr));
getsockname(listenfd, (struct sockaddr *)&servaddr, &len);
printf("bound name for client= %s\n", servaddr.sun_path);
for ( ; ; ) {
    sr=recvfrom(listenfd, buf, sizeof(buf), 0, (struct sockaddr *)&
        cliaddr, &len);
    write(1, buf, sr);
    write(1, "\n", 1);
}
return 0;
}
```

Unix Domain Datagram Client

```
int main(int argc, char **argv)
{
    int sockfd, cr, n;
    char buf[12];
    socklen_t len;
    struct sockaddr_un cliaddr;
    len = sizeof(cliaddr);
    sockfd = socket(AF_LOCAL, SOCK_DGRAM, 0);
    bzero(&cliaddr, sizeof(cliaddr));
    cliaddr.sun_family = AF_LOCAL;
    strcpy(cliaddr.sun_path, argv[1]);
    sendto(sockfd, "SOA---ITER", 10, 0, (struct sockaddr *)&cliaddr,
           len);
    write(1, "Me done\n", 8);
    return 0;
}
```

socketpair Function

The `socketpair` function creates two sockets that are then connected together. This function applies only to Unix domain sockets.

```
#include <sys/socket.h>
```

```
int
```

```
socketpair(int family, int type, int protocol, int sockfd[2]);
```

Returns: nonzero **if** OK, -1 on error

- ✎ The family must be **AF_LOCAL**
- ✎ The **type**, however, can be either **SOCK_STREAM** or **SOCK_DGRAM**.
- ✎ The **protocol** must be 0.
- ✎ The two socket descriptors that are created are returned as **sockfd[0]** and **sockfd[1]**.
- ✎ The two created sockets are unnamed; that is, there is no implicit **bind** involved.
- ✎ The result of `socketpair` with a type of **SOCK_STREAM** is called a **stream pipe**. It is similar to a regular Unix pipe (created by the `pipe` function), but a **stream pipe** is **full-duplex**; that is, both descriptors can be read and written.

Unix Domain Socket Uses: Passing Descriptors

When we think of passing an open descriptor from one process to another, we normally think of either

- ✎ A child sharing all the open descriptors with the parent after a call to **fork**
- ✎ All descriptors normally remaining open when **exec** is called

Current Unix systems provide a way to pass any open descriptor from one process to any other process. That is, there is no need for the processes to be related, such as a parent and its child. The technique requires us to first establish a Unix domain socket between the two processes and then use **sendmsg** to send a special message across the Unix domain socket. This message is handled specially by the kernel, passing the open descriptor from the sender to the receiver.

Receiving Sender Credentials

SELF-TRY.....

THANK YOU